

CS 224W Project Report

The Evolution of Wikipedia

Jiaji Hu, Haozhun Jin, Peng Qi
Group #48

1 Introduction

The evolution of online networks is a topic that has generated much interest in research. In the past, scientists have studied some dynamic properties of network evolution, choosing to focus on online social networks and proposing a variety of models with some degree of success. However, much is yet to be learned about how large networks grow and evolve.

As the largest online encyclopedic resource in human history, Wikipedia is an interesting example of an online network. Unlike social networks — the usual object of analysis, Wikipedia is a knowledge network, representing the structure of information instead of social relationships. Therefore, studying the dynamics of Wikipedia can be a refreshing addition to network evolution research.

In our project, we study the evolution of Wikipedia on both static and dynamic viewpoints. Using data that consists of Wikipedia's complete edit history, we take snapshots of Wikipedia at different timepoints and compute statistics to look for trends over time and gain insight into the network structure. By comparing snapshots taken between a short period of time, we single out individual edge creation processes and study them, focusing on the choice of destination for new edges. On our Wikipedia network, we compare and evaluate previous edge destination models such as the triangle closing model and models based on Preferential Attachment. From this, we go on to propose that new edges not only choose their destinations according to the degree of a destination node, but also according to the PageRank score of that destination. We evaluate the effectiveness of PageRank as a predictor of edge destination, comparing it to node degree. From our project, we hope to find out both results that reveal interesting information about Wikipedia in particular, and results that apply to network evolution in general.

2 Prior Work

In recent years, researchers have done a considerable amount of work on modeling the evolution of networks. In the process, both static and dynamic means of analysis have been used.

Static analysis consists of taking snapshots of the network at different stages of its development and computing statistical properties of interest for each snapshot to reach a conclusion to the process of network evolution. In [2], the authors took snapshots to analyze the structural properties of their network, focusing on the size of each connected component. In [4], the authors use snapshots to compute reciprocity of edges, arriving at conclusions about the structure of Wikipedia. However, using their model, the intermediate status of the graph in the process of generation does not match to any time point of the real graph. Only the final state of the generated graph matches the real-world graph in some characteristics.

On the other hand, research taking the dynamic view do not rely on snapshots, but rather studies the change in network on an edge-by-edge basis. In [3], the authors studied the exact edge arrival sequence to directly observe and model the fine-grain evolution process of networks. The paper considered dynamic analysis to be made up of three core processes: node arrival, edge initiation and edge destination. In particular, edge initiation and edge destination were closely studied, and a random-random triangle closing model was proposed as a good model for predicting edge destination.

The preferential attachment (PA) theory was proposed and evaluated in both [3] and [4]. In [3], PA was evaluated dynamically on social networks, while in [4], it was analyzed with a static view on Wikipedia data. This project will combine the two and assess preferential attachment with a dynamic point of view on the Wikipedia dataset, so as to further evaluate the validity of this theory.

We see that most papers in this field, including [3], [1] and [2] were written with social networks as the object of study. This raises the question whether their proposed models hold in other types of networks. This paper aims to explore some of these models on Wikipedia to answer this question.

3 Model

We model Wikipedia as an unweighted directed graph where each document is a node, and each edge $\langle u, v \rangle$ represents that there exists referential link from document u to document v . We model Wikipedia as unweighted for model simplicity, which means we choose to ignore the fact that some documents u might link to some other document v multiple times.

4 Data Collection and Processing

Wikipedia offers free copies of its contents for mirrors and research. In most of the use cases, a snapshot is enough. But for the purpose of our project, we need the edit history. Therefore we chose the *pages-meta-history* dataset, which contains the complete edit history. We used the August 2006 version of English Wikipedia from the official Wikimedia archives.

The dump is more than 720GB after decompression. To reduce the time cost of processing such data, for each article and for each revision, we extracted the article title, its list of outgoing links and time of revision. Next, we assigned each article an id to further minimize file sizes. This gave us a boiled down version of the information we needed. From there, we sorted all revisions along with the article id and ids of articles it links to by the revision time. After that, we were ready to generate snapshots of the Wikipedia graph over time.

Due to special characteristics of Wikipedia, we also did some additional preprocessing. For one, we removed all articles with a colon in its name. These articles were in special namespaces, serving as talk pages, user pages, images, or special administrative pages, etc. Such pages are not in our interest.

Also, redirections in Wikipedia may have a special effect on the graph. We decided to study this effect. To construct a redirection-free version of the graph, we removed all articles that redirected to another in the snapshots. After the removal, we moved all relevant edges to the page that it was redirected

to. We noticed a very small portion of redirections that turned into a cycle. Such pages were removed completely.

5 Snapshot Analysis

A snapshot is a static view of an evolving network at a certain point in time. After generating a snapshot, we can do static analysis of its properties. In our project, we looked at the following properties of each snapshot:

1. Number of nodes and edges.
2. Degree distributions, including in-degree and out-degree.
3. Diameter (sampling from a subset of nodes for efficiency).
4. Connected component sizes.
5. Average clustering coefficient.

On one hand, looking at these properties for each snapshot gives us an idea about the structure of the graph at that time. On the other hand, by comparing these statistic over many snapshots through time, we can start to look for trends as the network develops. We try to choose the interval between snapshots so that the snapshots are close enough to fully present the developmental stages of the network, while taking into account the limits of computation and memory. After generation and calculation, we visualize out results to show our findings more clearly. The results are as follows:

5.1 Degree Distribution

5.2 Degree Distribution

5.3 Degree Distribution

6 Dynamic Analysis

After we generate snapshots and analyze properties of the whole network, we move on to study the dynamics of the network, namely, the process that determines edge destination.

We study individual edge creation processes in the following way: Note that if we take two snapshots of

the network in a short period of time, the vast majority of nodes and edges will not have changed. If the interval is sufficiently short, there will only be a small amount of changes in the network, consisting of creations and deletions of nodes and edges. Since we are not interested in the deletion of nodes and edges, we will only look at the added nodes and edges. In this way, by comparing snapshots, we can single out individual node creations and edge creations. The process of comparing snapshots is shown below:




Figure 1: Comparing snapshots

This process is analogous to taking the derivative of the graph at a certain time. Now, for each observed edge creation, we can study it in a dynamic setting. As mentioned earlier, we want to study the selection of destination, in particular two aspects:

1. How far away the destination is from the source (Triangle closing model).
2. What are the properties of the destination node (Preferential Attachment model).

6.1 Triangle Closing

For the first item, we treat our network as undirected and run depth-limited BFS (Shortest Path would be computationally inefficient) from the source node to try to determine the number of hops between source and destination. Distance 1 means that the new edge is a reciprocal edge, distance 2 means that the new edge closes a triangle. We look at the proportion of reciprocal edges and triangle-closing edges. 2 shows the results.




Figure 2: Distance between source and destination of new edges

Then, we revert to treating the graph as directed to study the triangle closing case. Note that the same triangle closing case on an undirected graph can turn into four possible types of cases for a directed graph, as shown below:




Figure 3: Types of directed triangles

We look into which kind of triangles the new edges tend to close. Also, note that one single edge may close more than one type of triangle in a directed graph, and there are 15 combinations in all.




Figure 4: Combination frequency of triangles closed

6.2 Preferential Attachment

After that, we study Preferential Attachment(PA) in edge destination selection. Remember that in our class, the model was described as when a new node j joins the network, it creates a constant number of edges, where edge destination is proportional to the destination's degree.

$$P(j \rightarrow i) = \frac{d_i}{\sum_k d_k} \quad (1)$$

In [3], PA was studied as the bias in selection of edge destination based on the degree and age of the node. In the case of degree, the probability of a new edge e choosing a destination of degree d was computed as:

$$p_e(d) = \frac{\sum_t [e_t = (u, v) \wedge d_{t-1}(v) = d]}{\sum_t |u : d_{t-1}(u) = d|} \quad (2)$$

We will use the above equation to calculate $p_e(d)$ for edge creation in our graph. If the data was generated by the PA model given in (1), we would see that $p_e(d) \propto d$. If the attachment in our Wikipedia graph is close to linear, then the PA model by degree is further justified. If not, it would be an interesting phenomenon that would not be explained by PA by degree.

7 Model Proposal - Preferential Attachment by PageRank

After evaluating Preferential Attachment by degree on our Wikipedia graph, we go on to propose a PA model of our own. Recall that the original PA model arised from the notation that “the rich get richer”, with the degree of a node as a measure of importance. It is therefore natural to retain this notion of cumulative advantage, but use a different feature as the measure of importance.

One motivation for trying out another measure of importance arises out of observations from previous results. In [3], PA with in-degree was evaluated on some social networks. The findings were that $p_e(d) \propto d$ worked for some networks (for DELICIOUS, $p_e(d) \propto d^{0.9} \approx d$), but for some networks, links did not attach preferentially by degree. In the case of LINKEDIN, we see that $p_e(d) \propto d^{0.6}$ for low degree nodes and $p_e(d) \propto d^{1.2}$ for high degree nodes, suggesting that low degree and high degree nodes get sub-preferential and super-preferential treatment according to their degree, respectively. This leads one to contemplate the possibility that there exists a better measure m for importance than in-degree, so that $p_e(d) \propto m$ fits well to broad cases.

We propose that the PageRank score of a node may be a good alternative measure of importance. Recall that PageRank conveys the idea that the importance of a node not only depends on the number of links it has, but also depends on the importance of the nodes it is linked to. This appears to be a more robust notion of importance than just the in-degree of a node.

Therefore, we are interested in trying out Preferential Attachment by PageRank. This will be done as follows:

First, compare snapshots with short intervals to single out individual edge creations. Say we are studying the destination of newly created edge e .

Next, calculate the PageRank scores of nodes in the original graph using the equation below (with other implementation details):

$$r(j) = \sum_{i \rightarrow j} \beta \frac{r(i)}{d_i} + (1 - \beta) \frac{1}{n} \quad (3)$$

Then, just like how we evaluated PA by node degree, we calculate the probability of a new edge e

choosing a destination of PageRank pr as:

$$p_e(pr) = \frac{\sum_t [e_t = (u, v) \wedge r_{t-1}(v) = pr]}{\sum_t |u : r_{t-1}(u) = pr|} \quad (4)$$

(Note that we discretize the PageRank score of nodes so that the results are meaningful.)

Finally, we can plot $p_e(pr)$ against pr to see if there is a relation where $p_e(pr) \propto pr^\tau$, $\tau \approx 1$.

If PageRank does turn out to be a better measure of importance for Preferential Attachment, we can easily modify the generation model give by (1), substituting $\frac{d_i}{\sum_k d_k}$ with $\frac{r_i}{\sum_k r_k} = r_i$, which gives us

$$P(j \rightarrow i) = \frac{r_i}{\sum_k r_k} = r_i \quad (5)$$

In this generation model, after the insertion of each node, PageRank needs to be run to modify the new weights. Luckily, we can use values from the previous iteration (warm start) to greatly save computation.

8 Results

Up to now, we have finished a tough part of our work, which is processing the 720G dataset and generating our graph. After taking snapshots monthly for the network, we started to study the properties of the snap graphs.

From Fig. 5, it is evident that Wikipedia network grows almost exponentially over time in terms of both node number and edge number (especially after 2003). It can be further observed that the number of both Wikipedia’s nodes and its edges grow by about an order (10x) in about 22 months. Fig. 5 (right) shows that the number of nodes and edges in Wikipedia follows a rough linear ratio in their growth, indicating that Wikipedia is similar to a scale-free network.

A closer observation on the linear scale graph revealed a significant difference in node count before and after (marked with “(R)” in the legends, or the red lines/points) our redirection process, which is intuitively correct as it removes conceptually redundant nodes and edges. As for edge/node ratio, after processing redirection, the edge/node ratio still follows a roughly linear growth, which indicates that the ratio

of redirection nodes in the corpus is roughly constant.

Further analysis of Wikipedia’s evolution of diameter reveals an interesting phenomenon. As Fig. 6 shows, the diameter of graph increases first, and then decreases to a constant range and stabilizes. The initial large diameter might be caused by the disconnectedness of the graph, but the final stabilization further endorses our observation that Wikipedia is a scale-free network. It is not surprising that after processing redirections, the diameter of the graph shrinks by about 1, the common length of redirection chains.




Figure 6: Change of Diameter over Time

Besides those common benchmarks above, we also investigated the change of Wikipedia’s connected component sizes over time, which might lend some insight to the difference between the difference between typical social networks and knowledge networks like Wikipedia (more comparison to social networks will be required to draw a conclusion). To investigate the huge difference in SCC sizes before and after redirection processing, we analyzed the portion of redirection nodes that are inside the largest SCC. As it turned out, only about 40% of the removed redirection nodes are from the largest SCC, which suggest that redirection pages occur more in the in-component, out-component, and disjoint component of the largest SCC. This would require further investigation of the general structure of the Wikipedia network during the rest of our project.




Figure 7: Change of Connected Component Sizes

Finally, we also plotted the changes of maximum




Figure 5: Growth of Network over Time

in- and out-degrees over time, and briefly analyzed some interesting facts. Though these graphs do not provide valuable insights themselves (thus they will be replaced with the evolution of degree distributions in our final report), we did find “fun facts” analyzing these figures. For in-degree, the significant jump that took place by the end of 2002 was caused by page “2000”, as it turned out that people back then tend to link many occurrences of year 2000 to that page. As for large out-degree nodes, typically they are lists of entities, and the administrators remove some of the “unencyclopedia” lists from time to time, which reduces the maximum out-degree at certain time points.




Figure 8: Change of Max In- and Out-Degree over Time

9 Discussion and Next Steps

Looking back at the project proposal, we have made a huge amount of progress. From the project planning and theoretical point of view, we have set a focus on what we want to do, and brainstormed to propose our PA model with PageRank that is based on a current model (PA-degree model), aims to solve a problem that the current model does not explain (As explained in section 5), and actually has a chance to

do well. We feel that this is becoming an interesting topic to explore.

As of now, we have pretty much completed the static analysis of snapshots of the graph. Our next step would be to start looking into the dynamic edge creation process. Our two main tasks are as follows:

1. Use dynamic analysis to evaluate the triangle closing model and PA-degree model on our network.
2. Compute PageRank scores, evaluate our PA-PageRank model.

After finally finishing the data processing step, we have been enjoying good progress in an interesting project. We are looking forward to a good project.

References

- [1] Gueorgi Kossinets and Duncan J Watts. Empirical analysis of an evolving social network. *Science*, 311(5757):88–90, 2006.
- [2] Ravi Kumar, Jasmine Novak, and Andrew Tomkins. Structure and evolution of online social networks. In *Link Mining: Models, Algorithms, and Applications*, pages 337–357. Springer, 2010.
- [3] Jure Leskovec, Lars Backstrom, Ravi Kumar, and Andrew Tomkins. Microscopic evolution of social networks. In *Proceedings of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 462–470. ACM, 2008.
- [4] Vinko Zlatić and Hrvoje Štefanić. Model of wikipedia growth based on information exchange via reciprocal arcs. *EPL (Europhysics Letters)*, 93(5):58005, 2011.